

HP Vertica on Amazon Web Services Guide

HP Vertica Analytic Database

Software Version: 7.0.x



Document Release Date: 2/24/2014



Legal Notices

Warranty

The only warranties for HP products and services are set forth in the express warranty statements accompanying such products and services. Nothing herein should be construed as constituting an additional warranty. HP shall not be liable for technical or editorial errors or omissions contained herein.

The information contained herein is subject to change without notice.

Restricted Rights Legend

Confidential computer software. Valid license from HP required for possession, use or copying. Consistent with FAR 12.211 and 12.212, Commercial Computer Software, Computer Software Documentation, and Technical Data for Commercial Items are licensed to the U.S. Government under vendor's standard commercial license.

Copyright Notice

© Copyright 2006 - 2014 Hewlett-Packard Development Company, L.P.

Trademark Notices

Adobe® is a trademark of Adobe Systems Incorporated.

Microsoft® and Windows® are U.S. registered trademarks of Microsoft Corporation.

UNIX® is a registered trademark of The Open Group.

Contents

Contents	3
About Using HP Vertica on Amazon Web Services (AWS)	5
Important Notes on This Release of HP Vertica on AWS	5
Understanding the AWS Procedure Components	6
About the HP Vertica Instance Sizing	8
Instance Type Accommodates Large Clusters and High Query Concurrency	9
Elastic Block Storage Provides Optimal Cluster Performance	9
Ample Disk Space Allotted for the Operating System	9
Installing and Running HP Vertica on AWS: The Detailed Procedure	11
Creating a Placement Group	11
Creating a Key Pair	12
Creating a Virtual Private Cloud (VPC)	12
Creating and Assigning an Internet Gateway	13
Creating a Security Group	14
Adding Rules to a Security Group	14
Configuring and Launching an Instance	17
Assigning an Elastic IP	19
Connecting to an Instance	20
Forming a Cluster	21
After Your Cluster Is Up and Running	22
Using Management Console (MC) on AWS	23
Adding Nodes to a Running AWS Cluster	25
Launching New Instances to Add to an Existing Cluster	25
Including the New Instances As Nodes in Your Existing Cluster	25
Adding Nodes and Rebalancing the Database	25
Removing Nodes From a Running AWS Cluster	27
Preparing to Remove a Node	27
Removing Hosts From the Database	27
Removing Nodes From the Cluster	28

Stopping the AWS Instances (Optional)	28
Migrating to HP Vertica 7.0.x or later on AWS	31
Upgrading to the Latest HP Vertica AMI on AWS	33
Preparing to Upgrade Your AMI	33
Upgrading Your AMI	33
Troubleshooting: Checking Open Ports Manually	37
Using the Netcat (nc) Utility	37
Appendix: Installing and Running HP Vertica on AWS - QuickStart	41
We appreciate your feedback!	43

About Using HP Vertica on Amazon Web Services (AWS)

In the context of AWS, an HP Vertica node is an AWS instance that has been configured using the HP Vertica AMI. When you create your instances, you specify the HP Vertica AMI you want to install by using the `install_vertica` script.

This document covers the steps you should follow to create an HP Vertica cluster on AWS, from preparing your AWS environment and launching instances through combining those instances to create an HP Vertica cluster. You can choose to follow either the QuickStart procedure, or a more detailed procedure:

- The QuickStart procedure includes everything you need to know to set-up HP Vertica, but assumes you are familiar with the AWS Amazon Management Console. You can find the QuickStart procedure [near the end of this document](#).
- The detailed procedure adds steps to help walk you through setting up and combining instances to create a simple cluster.

Note that you run the `install_vertica` script at the end; the procedures in this document create your environment so that you can create a cluster using the `install_vertica` script.

Important Notes on This Release of HP Vertica on AWS

The instructions in this document apply to HP Vertica Version 7.0.x or later.

If you are running or upgrade to the latest HP Vertica Release:

- You can install the Management Console (MC) on your AWS cluster. For more information, see [Using Management Console \(MC\) on AWS](#).
- You can use the expanded HP Vertica Community Edition, which is the free version of the HP Vertica Analytics Platform. The free version remains limited to three nodes and up to 1 TB of data. Refer to *HP Vertica Community Edition* in the New Features Guide for more information on the enhanced Community Edition. Note that, if you are using Community Edition, the license is included within the AMI and is located here: `/opt/vertica/config/licensing/vertica_community_edition.license.key`

The HP Vertica AMI includes a pre-formatted 8-disk software RAID array mounted at `/vertica/data` (The location is the data and catalog path, which you use when you create a database.)

Understanding the AWS Procedure Components

The following illustration and table display and describe the AWS components that you must configure. These are preliminary steps in creating an HP Vertica cluster on Amazon AWS. You can skip this section and go right to the [Installing and Running HP Vertica on AWS: The Detailed Procedure](#).

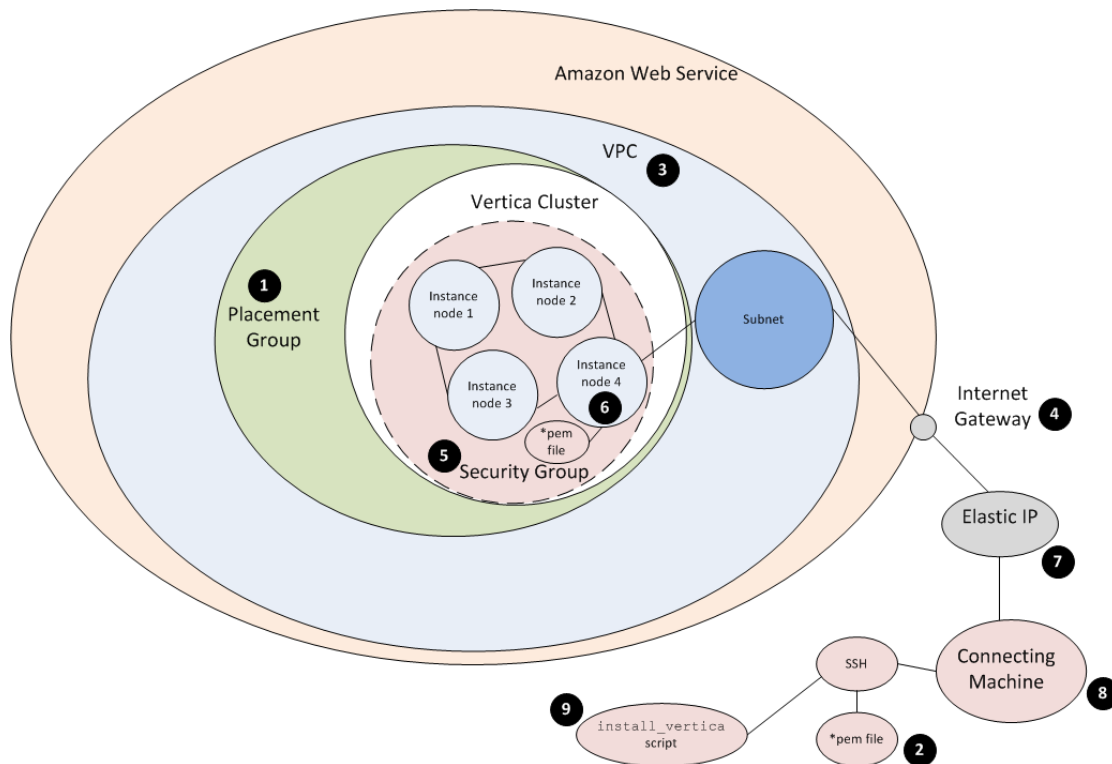


Illustration Reference and Step #	Component	Sample	Description
1	Placement Group	GroupVerticaP	You supply the name of a Placement Group when you create instances. You use a Placement Group to group instances together. A Placement Group for a cluster resides in one availability zone; your Placement Group cannot span zones. A Placement Group includes instances of the same type; you cannot mix types of instances in a Placement Group. You can choose one of two regions for a Placement Group; see Creating a Placement Group for information.

Illustration Reference and Step #	Component	Sample	Description
2	Key Pair	MyKey (MyKey.pem would be the resultant file.)	You need a Key Pair to access your instances using SSH. You create the Key Pair through the AWS interface, and you store a copy of your key (*.pem) file on your local machine. When you access an instance, you need to know the local path of your key, and you copy the key to your instance before you can run the <code>install_vertica</code> script.
3	VPC	Vpc-d6c18dbd Amazon assigns this name automatically.	You create a Virtual Private Cloud (VPC) on Amazon so that you can create a network of your EC2 instances. All instances within the VPC share the same network and security settings. A VPC uses the CIDR format for specifying a range of IP addresses.
4	Internet Gateway	Igw-d7c18dbc Amazon assigns and attaches an internet gateway automatically.	An internet gateway allows instances to access the internet. Note that, typically, a gateway is automatically assigned when you create a VPC. You can also create your own named internet gateway, and then attach that gateway to your VPC.
5	Security Group	"default" Amazon assigns a security group named "default" automatically, but you can create and name your own security group.	The security group includes firewall settings; its rules specify how traffic can get in and out of your instances. When you launch your instances, you choose a security group. If you use the default security group, you must add HP Vertica recommended rules to the group as you would if you created the group yourself.

Illustration Reference and Step #	Component	Sample	Description
6	Instances	<p>10.0.3.158 (all samples)</p> <p>10.0.3.159</p> <p>10.0.3.157</p> <p>10.0.3.160</p> <p>(Amazon assigns IP addresses. You need to list the addresses when you run the <code>install_vertica</code> script to form a cluster of the instances.)</p>	<p>An instance is a running version of the AMI. You first choose an AMI and specify the number of instances. You then launch those instances. Once the instances are running, you can access one or more of them through SSH.</p> <p>Notes:</p> <ul style="list-style-type: none"> Once IPs are assigned to your instances, they must not change in order for HP Vertica to continue working. The IP addresses are crucial once you have run the <code>install_vertica</code> script; if you change them or they become reassigned, your cluster breaks. When you stop one or more instances, your My Instances page may show blanks for the IP fields. However, by default Amazon retains the IPs and they should show up again in short order.
7	Elastic IP	107.23.104.78 (sample)	You associate an elastic IP to one of your instances so you can communicate with your cluster. An elastic IP is a static IP address that stays connected to your account until you explicitly release it.
8	Connecting	---	Once you have completed all of the procedures for setting up your cluster, you can connect to the instance attached to your elastic IP. You can use SSH. If connecting from a Windows machine, you can, for example, use Putty.
9	install_vertica script	---	Your instances are already grouped through Amazon and include the HP Vertica software, but you must then run the <code>install_vertica</code> script to combine instances through their IP addresses so that HP Vertica knows the instances are all part of the same cluster.

About the HP Vertica Instance Sizing

Sizing an instance involves estimating the hardware requirements for optimal HP Vertica performance in a typical scenario. HP Vertica's recommended pre-configured AMI is meant to apply to many typical customer environments; HP Vertica uses an AWS cluster-compute type instance in its AMI. HP Vertica provides its pre-configured AMI to ensure that CPU, RAM, storage I/O and

networking is allocated in a balanced manner, so that no particular resource bottleneck leaves other resources underutilized.

A *cluster-compute* instance includes a large amount of CPU and increased networking performance. Faster processors and substantially improved networking capabilities are important properties in scenarios such as operating clusters with a large node count loading very large volumes of data, or executing queries requiring on-the-fly data re-segmentation.

Note: You can read about Amazon cluster instances and their limitations in the [Amazon documentation](#).

Instance Type Accommodates Large Clusters and High Query Concurrency

The HP Vertica pre-configured AMI includes the AWS cc2.8xlarge instance type, allowing for very large clusters, higher query concurrency, continuous (trickle) loads, and high-volume periodic concurrent bulk loads.

Elastic Block Storage Provides Optimal Cluster Performance

EBS (Elastic Block Storage) is block-level storage used within Amazon EC2 instances. HP Vertica includes eight EBS volumes for each instance you create with the recommended HP Vertica AMI; you do not configure EBS volumes when you launch an instance. Every CPU within the instance receives I/O throughput independently. Prior to release 6.1, before HP Vertica provided a pre-configured AMI, under-configured EC2 storage was often the root cause of unsatisfactory AWS cluster performance.

The eight volumes of EBS storage are presented to the host instance as a software RAID0 device. HP Vertica has found that RAID0 device performance continues to increase up to the optimal number of eight volumes. Note that HA is provided through K-safety in HP Vertica. Refer to *High Availability and Recovery* in the Concepts Guide for information on K-safety.

HP Vertica uses a recommended size for each EBS volume of 146 GB. With eight volumes at 146 GB per volume, total volume size compares to a typical HP Vertica hardware recommendation of 1.2TB.

Ample Disk Space Allotted for the Operating System

The HP Vertica AMI includes 15 GB disk space for the OS, which HP Vertica has determined to be an adequate allotment.

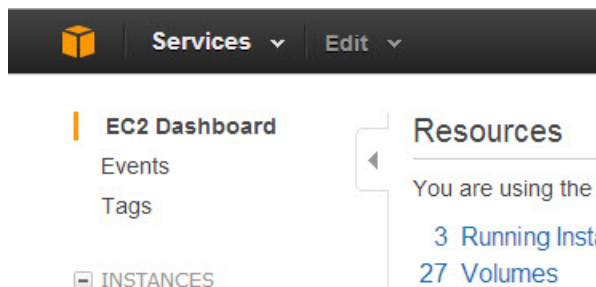
Installing and Running HP Vertica on AWS: The Detailed Procedure

Use these procedures to install and run HP Vertica on AWS. Note that this document mentions basic requirements only, by way of a simple example. Refer to the AWS documentation for more information on each individual parameter.

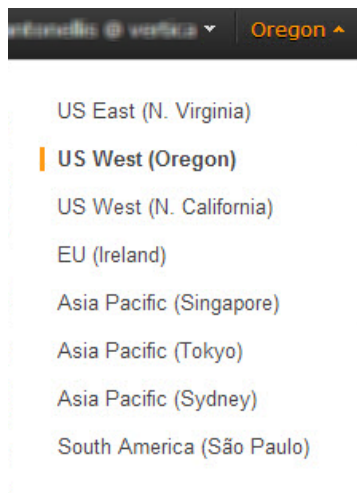
Creating a Placement Group

Perform the following steps to create a Placement Group. A Placement Group ensures that your nodes are properly co-located.

1. Log in to your Amazon EC2 Console Dashboard.



2. Select **Placement Groups**. The Placement Group screen appears.
3. Using the region dropdown on the top right of the Placement Group screen, choose one of two regions for your nodes. Two AWS regions support the HP Vertica AMI: **US West (Oregon)**, and **US East (N. Virginia)**.



4. Click **Create Placement Group**.
5. Name your placement group.
6. Click **Yes, Create**. Your group is created.

Creating a Key Pair

Perform the following steps to create a Key Pair.

1. Select **Key Pairs** from the Navigation panel.

Note: Depending upon which browser you are using, you may have to turn off the pop-up blocker in order to download the Key Pair.

2. Click **Create Key Pair**.
3. Name your Key Pair.
4. Click **Yes**.

The system displays a message letting you know that the Key Pair has been created.

5. Save your key pair. Ensure that you keep the *.pem file; you need it to logon to your instances.

Note that your Key Pair name now appears on the Key Pair list.

Creating a Virtual Private Cloud (VPC)

Perform the following steps to create a VPC.

1. From the AWS Management Console Home, navigate to the VPC console by selecting VPC.

Note: If the VPC Console Dashboard shows that VPCs already exist, you can select Your VPCs to note the names of the existing VPCs. As the VPC IDs are very similar, noting the names of the existing VPCs helps you later in identifying the new one you are about to create.

2. Click **Start VPC Wizard**.
3. From the wizard that displays, select **VPC with a Single Public Subnet Only**.
4. From the next screen, **Create an Amazon Virtual Private Cloud**, select **Edit Public Subnet**.

5. Change the Public Subnet as desired. HP Vertica recommends that you secure your network with an Access Control List (ACL) that is appropriate to your situation. The default ACL does not provide a high level of security.
6. Choose an Availability Zone.

Note: An HP Vertica cluster is operated within a single availability zone.

7. Click **Create VPC**. Amazon displays a message noting success.
8. Choose **Your VPCs** from the navigation pane, select your new VPC, and be sure that both **Enable DNS resolution** and **Enable DNS hostname support for instances launched in this VPC** are both checked.
9. Click **Close**.

Creating and Assigning an Internet Gateway

When you create a VPC, an internet gateway is automatically assigned to the VPC. You can use that gateway, or you can assign your own. To create and assign your own, perform the following steps. If using the default, continue with the next procedure, [Creating a Security Group](#).


1. From the navigation pane, choose **Internet Gateways**.
2. Click **Create Internet gateway**.
3. Choose **Yes, Create**.

The Internet Gateway screen appears. Note that the new gateway is not assigned to a VPC.

4. Click **Attach to VPC**.

5. Choose your VPC.
6. Click **Yes, Attach**.

Note that your gateway is now attached to your VPC.



Creating a Security Group

When you create a VPC, a default security group is created and automatically assigned. You can use that security group, or you can name and assign your own. To create and assign your own, perform the following steps. If using the default, continue with the next procedure, [Adding Rules to a Security Group](#). Note that you must add the HP Vertica rules as described in the next section.

To create and name your own security group, perform the following steps.

1. From the Navigation pane, select **Security Groups**.
2. Click **Create Security Group**.
3. Enter a name for the group and fill-in a description.
4. Choose your VPC.
5. Select **Yes, Create**.

Adding Rules to a Security Group

Perform the following steps to add rules to the security group you plan to use (whether you plan to use the default group or have created your own). This section includes procedures for adding rules from the **Inbound** tab and adding rule from the **Outbound** tab. You perform both procedures to establish your security group.

Perform the following to add rules from the **Inbound** tab:


1. Ensure that you have checked the box for the security group you just created, and that the **Inbound** tab is selected.

2. Add the HTTP rule.
 - a. Ensure that the **Inbound** tab is selected.
 - b. From the **Create a new rule** dropdown, choose **HTTP**.
 - c. Click **Add Rule**. The HTTP rule is added to your security group.
3. Add the Echo Reply rule.
 - a. From the **Create a new rule** dropdown, select **Custom ICMP rule**.
 - b. From the Type dropdown, select **Echo Reply**.
 - c. Click **Add Rule**. The Echo Reply rule is added to your security group.
4. Add the Traceroute rule.
 - a. Choose **Custom ICMP rule** once again.
 - b. Select **Traceroute**.
 - c. Click **Add Rule**. The Traceroute rule is added to your security group.
5. Add the SSH rule.
 - a. From the **Create a new rule** dropdown, choose **SSH**.
 - b. Click **Add Rule**. The SSH rule is added to your security group.
6. Add the HTTPS rule.
 - a. From the **Create a new rule** dropdown, choose **HTTPS**.
 - b. Click **Add Rule**. The HTTPS rule is added to your security group.
7. Add a port range to your security group.
 - a. From the **Create a new rule** dropdown, select **Custom TCP Rule**.
 - b. Under **Port range**, enter **4803-4805**.
 - c. Click **Add Rule**. The port range is added to your security group.
 - d. From the **Create a new rule** dropdown, select **Custom UDP Rule**.
 - e. Under **Port range**, enter **4803-4805**.
 - f. Click **Add Rule**. The port range is added to your security group.

8. Add individual ports.
 - a. Also under **Custom TCP rule**, enter the following ports under **Port Range**, one by one: **5433**, **5434**, **5444**, and **5450**. You enter the ports one by one to ensure the port assignments are sequential. HP Vertica uses these ports for internode communication.
 - b. Click **Add Rule** as you enter each number. Each port is added to your security group.
9. Click **Apply Rule Changes**.

Note: You must click **Apply Rule Changes** or your rules will not be applied to your security group. With the **Inbound** tab selected, your security group screen should look similar to the following.

1 Security Group selected


 **Security Group: vertica60**


Details **Inbound** Outbound Tags

Create a new rule: Custom TCP rule

Port range:
(e.g., 80 or 49152-65535)

Source:
(e.g., 192.168.2.0/24, sg-47ad482e, or 1234567890/default)

 Add Rule

 Apply Rule Changes

ICMP	Source
Echo Reply	0.0.0.0/0
Traceroute	0.0.0.0/0

TCP	Source
22 (SSH)	0.0.0.0/0
80 (HTTP)	0.0.0.0/0
443 (HTTPS)	0.0.0.0/0
4803 - 4805	0.0.0.0/0
5433 - 5434	0.0.0.0/0
5444	0.0.0.0/0

Perform the following to add rules from the **Outbound** tab.

Note: You want to ensure that all outbound traffic is allowed.

1. Select the **Outbound** tab.
 - a. Choose **All TCP** rule from the **Create a new rule** dropdown.
 - b. Click **Add Rule**. The All TCP rule is added to your security group.
2. Add the **All ICMP** rule.
 - a. Choose **All ICMP** rule from the **Create a new rule** dropdown.
 - b. Click **Add Rule**. The All ICMP rule is added to your security group.
3. Add the **ALL UDP** rule.

- a. Choose **ALL UDP** rule from the **Create a new rule** dropdown.
 - b. Click **Add Rule**. The ALL UDP rule is added to your security group.
4. Click **Apply Rule Changes**.

Note: You must click **Apply Rule Changes** or your rules will not be applied to your security group. With the **Outbound** tab selected, your screen should look similar to the following.

1 Security Group selected

Security Group: vertica60

Details Inbound **Outbound** Tags

Create a new rule: Custom TCP rule

Port range:
(e.g., 80 or 49152-65535)

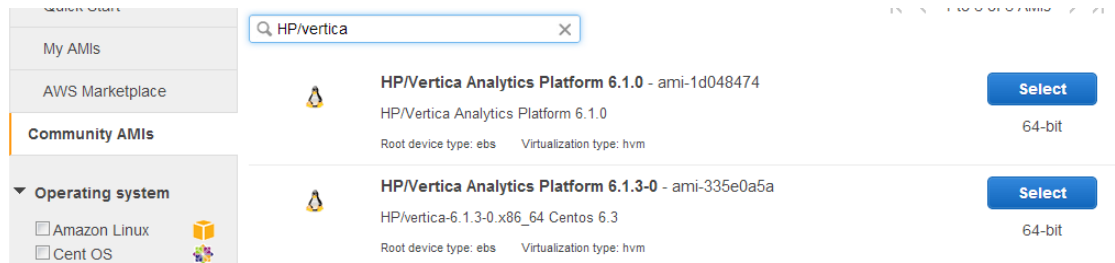
Destination: 0.0.0.0/0
(e.g., 192.168.2.0/24, sg-47ad482e, or 1234567890/default)

ALL	Port (Service)	Destination
ALL		0.0.0.0/0
ICMP	Port (Service)	Destination
ALL		0.0.0.0/0
TCP	Port (Service)	Destination
0 - 65535		0.0.0.0/0
UDP	Port (Service)	Destination
0 - 65535		0.0.0.0/0

Configuring and Launching an Instance

Perform the following steps to configure and launch the instances that will become your cluster in a later procedure. A basic EC2 instance (sans HP Vertica AMI) is similar to a traditional host. When you create an EC2 instance using an HP Vertica AMI, the instance includes the HP Vertica software and a standard recommended configuration to ease creation of your cluster. HP Vertica recommends that you use the HP Vertica AMI as is – without modification. The HP Vertica AMI acts as a template, allowing you to save configuration steps.

1. Navigate to the AWS Console, and select **EC2**.
2. From the EC2 Console Dashboard, click **Launch Instance**.
3. Choose the HP Vertica AMI.
 - a. Select the **Community AMIs** tab and enter "HP/Vertica" in the search field.
 - b. Choose your HP Vertica AMI. Click the **Select** button next to the AMI.



Note: The HP Vertica AMI has a preformatted 8 disk software RAID array mounted at `/vertica/data`. The location is the data and catalog path, which you use when creating a database.

4. From, **Choose an Instance Type**, select the **Compute optimized** tab and select **cc2.8xlarge**.
5. Click **Next:Configure Instance Details**.
 - a. Choose the number of instances you want to launch. For our example, we are entering **4**, which represents the number of nodes in our cluster.

Note: An HP Vertica cluster uses identically configured instances of the same type. You cannot mix instance types.

- b. From the **Network** drop-down, choose your VPC.
 - c. From the **Placement group** drop-down, choose the placement group you previously created.
6. Click through to **Configure Security Group**.

Note: As you click through, make no changes to the Add Storage screen, which shows eight EBS volumes. Do not change the settings on this screen (e.g., do not add or delete disks or the RAID array already set-up for you will be destroyed).

From **Select an existing security group**, choose the security group you previously created.

Note: Not all data centers support VPC. If you receive an error message that states "VPC is not currently supported...", choose a different subnet location (e.g., choose us-east-1c rather than us-east-1a).

7. Click **Review and Launch**. Review instance details.
8. Click **Launch**. Choose your key pair and check the box to acknowledge that you have access to the key file.

9. Click **Launch Instances**. You receive a message saying that your instances are launching.

You can click **View Instances** from the Launch Status page. Check that your instances are running (show green as their state).

Note: You can stop an instance by right-clicking and choosing **Stop**. Use Stop rather than Terminate; the latter will cause the instance to be lost.

Assigning an Elastic IP

Perform the following procedure to assign an Elastic IP address.

The elastic IP is an IP address that you attach to an instance; you communicate with your cluster through the instance that is attached to this IP address. An elastic IP is a static IP address that stays connected to your account until you explicitly release it.

1. From the Navigation menu, select **Elastic IPs**.
2. Click **Allocate New Address**. (Alternatively, you can select the checkbox for the address to associate with an Elastic IP.)
3. On the **Allocate New Address** screen, choose **VPC** from the dropdown.
4. Click **Yes, Allocate**.
5. Click **Associate Address**.
6. Choose one of the instances you created.
7. Click **Associate**.

Your instance is associated with your elastic IP.

Note the following on IP addresses:

- You can stop your instances, but you must ensure that the IP addresses assigned are not released.
- By default, Amazon keeps your IP addresses assigned unless you have selected "Allow Reassignment." To ensure your IP addresses will not be reassigned:
 - Right click on an instance and choose **Manage Private IP Addresses**.

- Ensure that **Allow Reassignment** is not checked. If it is checked, uncheck it.

The screenshot shows the 'Manage Private IP Addresses' dialog box. At the top, it says 'You can assign and unassign secondary private IP addresses on each network interface. Leave the address field blank and an available address will be assigned or enter an IP address that you want to assign.' Below this, it shows 'Instance: i-0eb0a273'. Underneath, there's a dropdown menu showing 'eth0: eni-42e0ac29 - Primary network interface - 10.0.3.0/24'. Below the dropdown, there are two input fields: '10.0.3.158' and '107.23.104.78', with 'Primary IP' written next to the second field. Below the input fields, there's a link that says 'Assign a secondary private address'. At the bottom, there's a checkbox labeled 'Allow reassignment' which is currently unchecked. At the bottom right, there are two buttons: 'Close' and 'Yes, Update'.

Connecting to an Instance

Perform the following procedure to connect to an instance within your VPC.

Note: Connect to the instance that is assigned to the Elastic IP.

1. Select **Instances** from the Navigation panel.
2. Select the instance that is attached to the Elastic IP.
3. Click **Connect**.
4. On **Connect to Your Instance**, choose from **A standalone SSH client** or **A Java SSH Client directly from my browser**.
 - If you have chosen **A Java SSH Client directly from my browser**:

Add the path to your private key in the field **Private key path**

Click **Launch SSH Client**.
 - If you have chosen **Connect with a standalone SSH client**, follow the steps required by your standalone SSH client.

Note that, if you connect to the instance from the Windows operating system, and plan to use Putty:

1. Convert your key file using PuTTYgen.
2. Connect with Putty or WinSCP (connect via the elastic IP), using your converted key (i.e., the *ppk file).
3. Move your key file (the *pem file) to the root dir using Putty or WinSCP.

Forming a Cluster

Use the `install_vertica` script to combine two or more individual instances and create a cluster.

The **My Instances** page lists the instances and their associated IP addresses. You need the IP addresses when you run the `install_vertica` script.

The following sample steps combine instances using the `install_vertica` script. Before you run `install_vertica`, be sure to stop any running databases on your nodes; the `install_vertica` script cannot complete successfully if any databases are running.

Note:

- When you issue `install_vertica` or `update_vertica` on an AMI, always use the `-T` (or `--point-to-point`) parameter to configure spread to use direct point-to-point communication between all HP Vertica nodes; this is a requirement for clusters on AWS. If you do not use the parameter, you receive an error telling you that you must use point-to-point communication on AWS.
- The `-d` parameter, if left out of your command, defaults to `/vertica/data` with AMIs.
- Password logons present a security risk on AWS. Include the parameter `--dba-user-password-disabled` so that the installer does not prompt for a password for the database user.

Refer to the Installation Guide for more information on the `install_vertica` script and all its options.

1. Locate your HP Vertica license.

Note: You need both the *pem file and your HP Vertica license to complete this procedure.

2. Copy files over to the instance.
 - a. Copy your *pem file (from where you saved it locally) onto the instance and place it in `/root`

Note: You must copy the file to the instance while logged on as root. Depending upon the procedure you use to copy the file, the permissions on the file may change. If permissions change, the `install_vertica` script will fail with a message similar to the following:

```
FATAL (19): Failed Login Validation 10.0.3.158, cannot resolve or connect to host as root.
```

If this happens, enter the following command to correct permissions on your `*pem` file:

```
chmod 600 /root/<name-of-pem>.pem
```

- b. Copy your HP Vertica license over to the instance and also place it in `/root`.
3. While connected to the instance, enter the following command to combine your instances into a cluster. The following is an example. Substitute the IP addresses for your instances and include your root `*pem` file name.

```
/opt/vertica/sbin/install_vertica -s  
10.0.3.100,10.0.3.101,10.0.3.98,10.0.3.99 -i ~/tst-key-pair.pem -L  
~/license.key -T --dba-user-password-disabled
```

Note: If you are using Community Edition, which limits you to three instances, you can simply specify `-L CE` with no license file.

For complete information on the `install_vertica` script and its parameters, see the Installation Guide, specifically the section, *About the install_vertica Script*.

After Your Cluster Is Up and Running

Stop or reboot instances using the Amazon AWS console, but you must stop the database before doing so. Once your cluster is up and running, if you need to stop or reboot:

1. Stop the database.
2. Stop or reboot one or more instances.

Caution: If you stop or reboot an instance (or the cluster) without shutting the database down first, disk or database corruption could result. Shutting the database down first ensures that HP Vertica is not in the process of writing to disk when you shutdown. Refer to the HP Vertica Administrator's Guide for information on stopping a database.

Once your cluster is configured and running:

1. Create a database. When HP Vertica was installed for you, an HP Vertica database administrator was created, `dbadmin`. You can use this pre-created `dbadmin` user to create and start a database. Refer to the HP Vertica Installation Guide for information on the `dbadmin` administrator.

2. Configure a database. Refer to the HP Vertica Administrator's Guide for information on configuring a database.
3. Refer to the full documentation set for HP Vertica for other tasks.

Using Management Console (MC) on AWS

MC is a database management tool that provides a way for you to view and manage aspects of your HP Vertica cluster. If you are running HP Vertica Release 6.1.2 or later, you can install and run MC.

This release of MC on AWS includes restrictions.

- You cannot create a cluster on AWS using the MC. You cannot import a cluster into AWS using the MC.
- You cannot monitor an AWS cluster using MC on a node that is outside of your AWS cluster. You must install MC on an instance within the AWS cluster itself.

Note: Each version of HP Vertica Management Console (MC) is compatible only with the matching version of the HP Vertica server. Version numbers must match to three digits; for example, HP Vertica 6.1.2 server is supported with HP Vertica 6.1.2 MC only. This is a general MC requirement and is not specific to MC on AWS.

What follows is a reading path for learning more about MC:

- For an overview of MC, where you can get it, and what you can do with it, refer to the Concepts Guide, specifically, *Management Console*.
- For information on installing and configuring MC, refer to the Installation Guide, specifically, *Installing and Configuring Management Console (MC)*.
- For information on the differences between what you can do with MC versus what you can do with the Administration Tools, refer to the Administration Guide, specifically the section, *Administration Tools and Management Console*.
- For information on creating a database using MC, refer to the Getting Started Guide, specifically the section, *Create the Example Database Using Management Console*.

Keep the following in mind concerning user accounts and the MC.

- When you first configure MC, during the configuration process you create an MC super user (a Linux account). Issuing a Factory Reset on the MC does not create a new MC super user, nor does it delete the existing MC super user. When initializing after a Factory Reset, you must logon using the original MC super user account.

For information on setting MC to its original state (Factory Reset), and why you might implement a Factory Reset, refer to the Administration Guide, specifically the section, [Resetting MC to Pre-configured state](#).

- Note that, once MC is configured, you can add users that are specific to MC. Users created through the MC interface are MC specific. When you subsequently change a password through the MC, you only change the password for the specific MC user. Passwords external to MC (i.e., system Linux users and HP Vertica database passwords) remain unchanged.

For information on MC users, refer to the Administration Guide, specifically the sections, [Creating an MC User](#) and [MC configuration privileges](#).

Adding Nodes to a Running AWS Cluster

Use these procedures to add instances/nodes to an AWS cluster. The procedures assume you have an AWS cluster up and running and have most-likely accomplished each of the following.

- Created a database.
- Defined a database schema.
- Loaded data.
- Run the database designer.
- Connected to your database.

Launching New Instances to Add to an Existing Cluster

Perform the procedure in [Configuring and Launching an Instance](#) to create new instances that you then will add to your existing cluster. Be sure to choose the same details you chose when you created the original instances (e.g., VPC and Placement group).

Including the New Instances As Nodes in Your Existing Cluster

Note that the **Instances** page lists the instances and their associated IP addresses. You need the IP addresses when you run the `install_vertica` script.

Perform the following to add the new instances as nodes to your existing cluster.

1. Connect to the instance that is assigned to the Elastic IP. See [Connecting to an Instance](#) if you need more information.
2. Enter the following command to add the new instances as nodes to your cluster. The following is an example. Substitute the IP addresses for your instances and include your root `*pem` file name. Your instances are added to your existing cluster.

```
/opt/vertica/sbin/install_vertica -A 10.0.3.110,10.0.3.111,10.0.3.112 -i  
~/tst-key-pair.pem -T
```

Adding Nodes and Rebalancing the Database

Once you have added the new instances to your existing cluster, you add them as nodes to your cluster, and then rebalance the database.

Follow the procedure given in the Administration Guide, [Adding Nodes to a Database](#).

Removing Nodes From a Running AWS Cluster

Use these procedures to remove instances/nodes from an AWS cluster.

Preparing to Remove a Node

Removing one or more nodes consists of the following general steps. The first two steps, backing up a database and lowering the k-safety before node removal, are prerequisites for the subsequent steps.

1. Back up the Database. See the section, *Creating Full and Incremental Snapshots (vbr.py)* in the Administrator's Guide.

HP recommends that you back up the database before performing this significant operation because it entails creating new projections, deleting old projections, and reloading data.

2. Lower the K-safety of your database if the cluster will not be large enough to support its current level of K-safety after you remove nodes. See the section, *Lowering the K-safety Level to Allow for Node Removal* in the Administrator's Guide.

Note: You cannot remove nodes if your cluster would not have the minimum number of nodes required to maintain your database's current **K-safety** level (3 nodes for a database with a K-safety level of 1, and 5 nodes for a K-safety level of 2). To remove the node or nodes from the database, you first must reduce the K-safety level of your database.

3. [Remove the hosts](#) from the database.
4. Remove the nodes from the cluster if they are not used by any other databases.
5. Optionally, [stop the instances](#) within AWS that are no longer included in the cluster.

Removing Hosts From the Database

Before performing the procedure in this section, you must have completed the tasks referenced in *Preparing to Remove a Node*. The following procedure assumes that you have both backed up your database and lowered the k-safety.

Note: Do not stop the database.

Perform the following to remove a host from the database.

1. While logged on as dbadmin, launch Administration Tools.

```
$ /opt/vertica/bin/admintools
```

Note: Do not remove the host that is attached to your EIP.

2. From the **Main Menu**, select **Advanced Tools Menu**.
3. From **Advanced Menu**, select Cluster Management. Select **OK**.
4. From **Cluster Management**, select **Remove Host(s)**. Select **OK**.
5. From **Select Database**, choose the database from which you plan to remove hosts. Select **OK**.
6. Select the host(s) to remove. Select **OK**.
7. Click **Yes** to confirm removal of the hosts.

Note: Enter a password if necessary. Leave blank if there is no password.

8. Select **OK**. The system displays a message letting you know that the hosts have been removed. Automatic re-balancing also occurs.
9. Select **OK** to confirm. Administration Tools brings you back to the **Cluster Management** menu.

Removing Nodes From the Cluster

To remove nodes from a cluster, run the `install_vertica` script, specifying the IP addresses of the nodes you are removing and the location and name of your root `*pem` file. (The following example removes only one node from the cluster.)

```
/opt/vertica/sbin/install_vertica -R 10.0.3.43 -i ~/tst-key-pair.pem -T
```

Stopping the AWS Instances (Optional)

Once you have removed one or more nodes from your cluster, to save costs associated with running instances, you can choose to stop or terminate the AWS instances that were previously part of your cluster. This step is optional because, once you have removed the node from your HP Vertica cluster, HP Vertica no longer sees the instance/node as part of the cluster even though it is still running within AWS.

To stop an instance in AWS:

1. On AWS, navigate to your **Instances** page.
2. Right-click on the instance, and choose **Stop**.

Migrating to HP Vertica 7.0.x or later on AWS

If you had an HP Vertica installation running on AWS prior to Release 6.1.x, you can migrate to HP Vertica 7.0.x or later using a new preconfigured AMI.

For more information, see the Solutions tab of the [myVertica portal](#).

Upgrading to the Latest HP Vertica AMI on AWS

Use these procedures for upgrading to the latest HP Vertica AMI. The procedures assume that you have a 6.1.x or later cluster successfully configured and running on AWS. If you are setting up an HP Vertica cluster on AWS for the first time, follow the [detailed procedure for installing and running HP Vertica on AWS](#).

Note: Both `install_vertica` and `update_vertica` use the same parameters.

Preparing to Upgrade Your AMI

Perform this procedure to prepare for the upgrade to the latest HP Vertica AMI.

1. Back up your existing database. See *Backing Up and Restoring the Database* in the Administrator's Guide.
2. Download the HP Vertica install package. See *Download and Install the HP Vertica Install Package* in the Installation Guide.

Upgrading Your AMI

Perform this procedure to upgrade to the latest HP Vertica AMI.

Note: All output is for illustrative purposes only and may change slightly from release to release.

1. Stop any running databases. You can use the Admin Tools dialog or a command line. A command line example follows.

```
[root@ip-10-0-10-117 ~]# sudo -u dbadmin /opt/vertica/bin/admintools
-t stop_db -d testdb
Info: no password specified, using none
Issuing shutdown command to database
Database testdb stopped successfully
```

2. Upgrade the rpm. A sample rpm command and output follows.

```
[root@ip-10-0-10-117 ~]# rpm --upgrade vertica-6.1.x.x86_64.RHEL5.rpm
Shutting down vertica services for upgrade
```

```
...
Shutting down vertica agent daemon
Stopping vertica agent:
Vertica Analytic Database V6.1.1-20130107 successfully installed on host ip-10-0-10-8
4
-----
Important Information
-----
If you are upgrading from a previous version,
you must backup your database before continuing with this install.
After restarting your database, you will be unable to revert to a previous
version of the software.
-----
To download the latest Vertica documentation in zip or tar format please visit http://
/www.vertica.com/documentation
To complete installation and configuration of the cluster,
run: /opt/vertica/sbin/install_vertica
```

3. Run the `update_vertica` script. You must use the `-i` option to specify location of your private-key file. Note that you cannot run the `install_vertica` script with the `sudo` command when you specify the `-i` parameter. An example follows.

Note: Refer to *About the `install_vertica` script* in the Installation Guide for more information on the `update_vertica` parameters. The `install_vertica` and `update_vertica` scripts use the same parameters.

```
[root@ip-10-0-10-117 ~]# /opt/vertica/sbin/update_vertica -u dbadmin
-r ~/vertica-6.1.x.x86_64.RHEL5.rpm -L ~/license.key.unlimited
-i ~/tst-key-pair.pem -T
Vertica Analytic Database 6.1.x Installation Tool
Starting installation tasks...
Getting system information for cluster (this may take a while)...
backing up admintools.conf on 10.0.10.118
backing up admintools.conf on 10.0.10.117
Installing rpm on 1 hosts...
installing node.... 10.0.10.118
NTP service is not running on the hosts: ['10.0.10.118', '10.0.10.117']
Vertica recommends that you keep the system clock synchronized using
NTP or some other time synchronization mechanism to keep all hosts
synchronized. Time variances can cause (inconsistent) query results
when using Date/Time Functions. For instructions, see:
* http://kbase.redhat.com/faq/FAQ_43_755.shtm
* http://kbase.redhat.com/faq/FAQ_43_2790.shtm
Info: the package 'pstack' is useful during troubleshooting.
Vertica recommends this package is installed.
Checking/fixing OS parameters....
Setting vm.min_free_kbytes to 31631 ...
Detected cpufreq module loaded on 10.0.10.118
Detected cpufreq module loaded on 10.0.10.117
CPU frequency scaling is enabled.
```

```
This may adversely affect the performance of your database.
Vertica recommends that cpu frequency scaling be turned off or set to 'performance'
Creating/Checking Vertica DBA group
Creating/Checking Vertica DBA user
Installing/Repairing SSH keys for dbadmin
Modifying existing Vertica DBA user
Creating Vertica Data Directory...
Skipping N-way network test (not recommended)
Updating spread configuration...
Verifying spread configuration on whole cluster.
Error Monitor 0 errors 4 warnings
Installation completed with warnings.
Installation complete.
To create a database:
1. Logout and login as dbadmin.**
2. Run /opt/vertica/bin/adminTools as dbadmin
3. Select Create Database from the Configuration Menu
** The installation modified the group privileges for dbadmin.
If you used sudo to install vertica as dbadmin, you will
need to logout and login again before the privileges are applied.
```

4. Start the database.
5. Check your HP Vertica version to ensure that upgrade was successful.

Troubleshooting: Checking Open Ports Manually

You originally configured your security group through the AWS interface. Once your cluster is up and running, you can check ports manually through the command line using the netcat (nc) utility. What follows is an example using the utility to check ports.

Before performing the procedure, choose the private IP addresses of two nodes in your cluster.

The examples given below use nodes with the private IPs:

```
10.0.11.60 10.0.11.61
```

Using the Netcat (nc) Utility

After installing the nc utility on both nodes, you can issue commands to check the ports on one node from another node.

1. Logon to the node that is attached to your elastic IP.

Note: The node must have internet access to perform the next step. If the node does not have internet access, you can copy the rpm to the node rather than download it directly from the internet.

2. Issue the following command to download the netcat (nc) utility.

```
yum install nc
```

3. Type y and press **Enter** when asked to confirm the download. A full example follows.

```
[root@ip-10-0-11-60 ~]# yum install nc
Loaded plugins: fastestmirror, security
Determining fastest mirrors
* base: mirrors.lga7.us.voxel.net
* extras: mirror.dattobackup.com
* updates: centos.mirror.netriplex.com

base                                | 3.7 kB    00:00
base/primary_db                    | 4.4 MB    00:00
extras                             | 3.5 kB    00:00
extras/primary_db                  | 19 kB     00:00
updates                            | 3.4 kB    00:00
updates/primary_db                 | 3.1 MB    00:00
Setting up Install Process
```

```
Resolving Dependencies
--> Running transaction check
---> Package nc.x86_64 0:1.84-22.el6 will be installed
--> Finished Dependency Resolution

Dependencies Resolved

=====
Package           Arch             Version           Repository        Size
=====
Installing:
nc                x86_64           1.84-22.el6       base              57 k
Transaction Summary
=====
Install            1 Package(s)

Total download size: 57 k
Installed size: 109 k
Is this ok [y/N]: y
Downloading Packages:
nc-1.84-22.el6.x86_64.rpm                | 57 kB      00:00
Running rpm_check_debug
Running Transaction Test
Transaction Test Succeeded
Running Transaction
Warning: RPMDB altered outside of yum.

Installing : nc-1.84-22.el6.x86_64
1/1
Verifying   : nc-1.84-22.el6.x86_64
1/1

Installed:

nc.x86_64 0:1.84-22.el6

Complete!
[root@ip-10-0-11-60 ~]#
```

4. Logon to the other node and install nc.
5. Stop spread on both nodes.

```
[root@ip-10-0-11-60 ~]# sudo /etc/init.d/spreadd stop
[root@ip-10-0-11-61 ~]# sudo /etc/init.d/spreadd stop
```

6. To check a TCP port:
 - a. Put one node in listen mode and specify the port. In the following sample, we're putting IP 10.0.11.60 into listen mode for port 4804.

```
[root@ip-10-0-11-60 ~]# nc -l 4804
```

- b. From the other node, run nc specifying the IP address of the node you just put in listen mode, and the same port number.

```
[root@ip-10-0-11-61 ~]# nc 10.0.11.60 4804
```

- c. Enter sample text from either node and it should show up on the other. To cancel once you have checked a port, enter **Ctrl+C**.

Note: To check a UDP port, use the same nc commands with the `-u` option.

```
[root@ip-10-0-11-60 ~]# nc -u -l 4804  
[root@ip-10-0-11-61 ~]# nc -u 10.0.11.60 4804
```


Appendix: Installing and Running HP Vertica on AWS - QuickStart

These steps cover specifics you need to know while working with the Amazon AWS interface. This procedure assumes you have a working knowledge of the AWS Management Console.

Note: This section covers all of the steps given in the full procedure, but is abbreviated for those familiar with the AWS Management Console. If you are not familiar with the AWS Management Console, please follow the [full procedure](#).

Steps from the AWS Management Console.

1. Choose a region, and create and name a Placement Group.
2. Create and name a Key Pair.
3. Create a VPC.
 - a. Edit the public subnet; change according to your planned system set-up.
 - b. Choose an availability zone.
4. Create an Internet Gateway and assign it to your VPC (or use the default gateway).
5. Create a Security Group for your VPC (or use the default security group, but add rules).

Inbound rules to add:

- HTTP
- Custom ICMP Rule: Echo Reply
- Custom ICMP Rule: Traceroute
- SSH
- HTTPS
- Custom TCP Rule: 4803-4805
- Custom UDP Rule: 4803-4805
- Custom TCP Rule: (Port Range) 5433
- Custom TCP Rule: (Port Range) 5434

- Custom TCP Rule: (Port Range) 5434
- Custom TCP Rule: (Port Range) 5450

Outbound rules to add:

- All TCP: (Destination) 0.0.0.0/0
- All ICMP: (Destination) 0.0.0.0/0

6. Create instances.
 - a. Select **Launch Instance**.
 - b. From Community AMIs, choose an HP Vertica AMI.
 - c. Select the **Compute optimized** tab and select **cc2.8xlarge**.
 - d. Choose the number of instances, network, and placement group.
 - e. Select a security group.
 - f. Click **Launch** and choose your key pair.
 - g. Click **Launch Instances**.
7. Assign an elastic IP to an instance.
8. Connect to the instance that is attached to the elastic IP.
9. Run the `install_vertica` script after placing the `*pem` file and the HP Vertica license in `/root`.

We appreciate your feedback!

If you have comments about this document, you can [contact the documentation team](#) by email. If an email client is configured on this system, click the link above and an email window opens with the following information in the subject line:

Feedback on HP Vertica on Amazon Web Services Guide (Vertica Analytic Database 7.0.x)

Just add your feedback to the email and click send.

If no email client is available, copy the information above to a new message in a web mail client, and send your feedback to vertica-docfeedback@hp.com.